

Search...

[Data Science](#)[Data Science Projects](#)[Data Analysis](#)[Data Visualization](#)[Machine Learning](#)[ML Projects](#)[I](#)

Image Classification using ResNet

Last Updated : 13 Feb, 2025

This article will walk you through the steps to implement it for image classification using Python and TensorFlow/Keras.

[Image classification](#) classifies an image into one of several predefined categories. **ResNet (Residual Networks)**, which introduced the concept of **residual connections** to address the vanishing gradient problem in very deep neural networks.

Here are the key reasons to use **ResNet** for image classification:

- **Enables Deeper Networks:** ResNet makes it possible to train networks with hundreds or even thousands of layers without performance degradation.
- **Improved Performance:** By using residual learning, ResNet achieves better accuracy in tasks like image classification.
- **Better Generalization:** The architecture helps avoid overfitting, improving model performance on unseen data.

Image Classification Using ResNet on CIFAR-10

Here's a step-by-step guide to implement image classification using the CIFAR-10 dataset and ResNet50 in TensorFlow:

1. Import Libraries

We begin by importing the necessary libraries from TensorFlow and Keras:

```
import tensorflow as tf
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.preprocessing.image import
ImageDataGenerator
```

```
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.layers import Dense, GlobalAveragePooli
7 from tensorflow.keras.optimizers import Adam
8 from tensorflow.keras.utils import to_categorical
```

2. Load and Preprocess the CIFAR-10 Dataset

We load the CIFAR-10 dataset using `tensorflow.keras.datasets.cifar10`. Then, we normalize the pixel values of the images (by dividing by 255) to scale them to a range of 0 to 1. Lastly, we one-hot encode the labels to match the output format for categorical classification.

```
1 # Load CIFAR-10 dataset
2 (x_train, y_train), (x_test, y_test) = cifar10.load_data()
3
4 # Preprocess the data
5 x_train = x_train.astype('float32') / 255.0
6 x_test = x_test.astype('float32') / 255.0
7
8 # One-hot encode the labels
9 y_train = to_categorical(y_train, 10)
10 y_test = to_categorical(y_test, 10)
```

3. Load ResNet50 Pre-trained on ImageNet

We use ResNet50, pre-trained on the ImageNet dataset. The `include_top=False` parameter ensures that the fully connected layers (the classification head) are not included, so we can add our custom layers.

```
1 base_model = ResNet50(weights='imagenet',
2                        include_top=False,
3                        input_shape=(32, 32, 3))
4
5 # Freeze the base model
6 base_model.trainable = False
```

4. Build the Classification Model

We now build the model using the pre-trained **ResNet50** as a base. We add